

Understanding Mutexes



David Flynn

SOFTWARE ENGINEER, FLYNN IT LTD



In the Last Module...



Saw how to share memory safely and avoid data races

- Using volatile
- Publishing immutable objects

Is this the only worry we might have?

Race Conditions



Race conditions lead to:

- Inconsistent data
- Invariants broken

Volatile cannot help

Publishing objects helps if only one thread may write

Mutual Exclusion

**Allows one
thread to access
a critical section
at once**

**Built in
'synchronized'
gives mutual
exclusion**

**Prevents
concurrent
access to
resources**



A Joke...



Two threads walk into a bar...



Thread One


```
enterBarAndOrderDrinks()  
roundBeingBought == false // 1  
roundBeingBought = true // 2  
buyDrinks()
```

Thread Two

```
enterBarAndOrderDrinks()  
roundBeingBought == true // 1  
waitForDrink()
```



Race Conditions




Incorrect behaviour caused by threads interleaving and executing code in an unintended order

Code usually works but occasionally goes wrong

- Code is not thread safe

Critical Section



One or more parts of the code which may not be accessed by more than one thread at a time

Critical sections may span more than one part of the code

Several critical sections may exist in the same program

Increment and Decrement

**Not atomic if can be
divided into smaller
pieces**

**Increment is not atomic:
Load, add, store**



Thread One

i starts
at 0

```
load i  
i = i + 1  
store i
```

i++

i is now 1

i is now 1,
not 2 as
expected

Thread Two

```
load i  
i = i + 1  
store i
```

i++

i is now 1



AtomicInteger




Provides safe pre and post increment and decrement

Cannot be divided into smaller parts
- So no interleaving

```
// static volatile int counter;  
static final AtomicInteger counter = new AtomicInteger();  
  
public void incrementAndDecrement()  
{  
    counter.getAndIncrement(); // instead of counter++;  
    counter.incrementAndGet(); // instead of ++counter;  
    counter.getAndDecrement(); // instead of counter--;  
    counter.decrementAndGet(); // instead of --counter;  
}
```



Mutex




Short for mutual exclusion object

**Prevents mutual access like a
key to a bathroom**

- Others wanting access have to wait

Demonstrating Race Conditions



Problems occur occasionally

- Impractical to manually run over and over

Could try:

- Putting code in a loop
- Using sleeps
- Using a `CountDownLatch`

```
if (!roundBeingBought) // 1
{
    roundBeingBought = true; // 2
    buyDrinks();
}
else
{
    waitForDrink();
}
```

**Critical
Section**



Obtaining Mutual Exclusion

Use synchronized
keyword

Implemented with
monitor locks (aka
intrinsic locks)



Implementation in Java



Objects are associated with intrinsic lock (a monitor)

Threads entering the monitor are put in the entry set

If no thread owns the monitor, one thread is chosen from the set

Thread acquires monitor and enters the critical section

Implementation in Java



**If another thread owns the monitor,
thread moves to Blocked state**

**Thread leaves gives up ownership
releasing the monitor**

Thread then exits the monitor

**Another thread is chosen from the entry
set moving from Blocked to Runnable**

Selection from Entry Set



No guarantee which thread chosen

**Possible under load that some threads
may never get chosen**

- No guarantee of fairness

Knowing if a Lock Is Being Held



We don't actually get a mutex object

- Can check `Thread.holdsLock(object)` to check if we own the monitor

Reentrancy



Monitors are reentrant

- Means we can acquire one if we own it without blocking
- Blocking would cause the thread to lock up (deadlock)
- Problem with some non-Java implementations

Probably should only check `holdsLock` to debug

Monitors and Objects


Monitor is associated with the object, not the reference

Should prevent reference changing (use final)

Same object for all threads that share critical section



Creating an Object to Use for Synchronization



```
private static final Object mutex = new Object();
```

Object is smallest to create so we're using that

Hide from other parts of the code

Should synchronize on different objects for different critical sections

- Using Boolean is bad as there are only two instances

```
synchronized(mutex)
{
    if (!roundBeingBought) // 1
    {
        roundBeingBought = true;
    }
}
```


```
// did we set the flag or not and thus should buy drinks?
```




```
boolean shouldBuyDrinks = false;
synchronized(mutex)
{
    if (!roundBeingBought) // 1
    {
        roundBeingBought = true; // 2
        shouldBuyDrinks = true;
    }
}
// now just need to test shouldBuyDrinks
```



AtomicBoolean



Using mutual exclusion for
'test-and-set' heavyweight

AtomicBoolean is better for
implementing 'test-and-set'

Use method: `compareAndSet(boolean expect,
boolean update)`

- If the value is 'expect', it is set to 'update'
- Returns true if the set took place

```
private static final AtomicBoolean roundBeingBought =  
                                                    new AtomicBoolean();  
...  
if (roundBeingBought.compareAndSet(false, true)) {  
    buyDrinks();  
} else {  
    waitForDrink();  
}
```



Thread One

```
roundBeingBought = false  
roundBeingBought == false  
buyDrinks(visit)  
numRoundsBought++
```

Thread Two

```
roundBeingBought = false  
roundBeingBought == false  
buyDrinks(visit)  
numRoundsBought++
```



CountDownLatch



Only good for a single use

- Need to create when needed
- Or
- In advance

CyclicBarrier



Like a reusable CountdownLatch
- But needs to be reset



Phaser



CyclicBarrier with multiple phases [avoids needing to be reset]

When all threads reach the barrier

- Phase is incremented
- Threads may proceed

Thread One

```
roundBeingBought = false
```

```
arriveAndAwaitAdvance()
```

```
roundBeingBought == false } (atomic)  
roundBeingBought = true   ]  
buyDrinks()  
numRoundsBeingBought++
```

Thread Two

```
roundBeingBought = false
```

```
arriveAndAwaitAdvance()
```

```
roundBeingBought == true  
waitForDrink()
```



Thread One

```
roundBeingBought = false
```

```
arriveAndAwaitAdvance()
```

```
roundBeingBought == false } (atomic)  
roundBeingBought = true  
buyDrinks()
```

```
roundBeingBought = false  
arriveAndAwaitAdvance()
```

→
Next
round

Thread Two

```
roundBeingBought = false
```

```
arriveAndAwaitAdvance()
```

```
// also buys drinks
```




Synchronized Using the Same Object



If using same object on multiple blocks

- Only one thread can execute in any of them at once
- They become part of the same critical section

Synchronized Keyword on Method Definitions



For instance - methods, object that 'this' refers to will be synchronized on

- All object's synchronized methods become a single critical section
- Controlled by the same monitor

Benefits:

- Visible
- Documentable
- No statements can be inserted before

Method Signatures and Synchronized



Synchronized does not form part of the method signature

Two forms (one synchronized, one not) not possible. Instead:

- Two versions of the class (e.g. StringBuffer & StringBuilder)
- Two differently named or parameterized methods

Synchronized version can call non-synchronized version

Hashtable



Thread safe HashMap with synchronized methods

- Doesn't actually call HashMap

Get and put can't be used simultaneously

- Get doesn't modify
- Synchronized prevents multiple readers
- May cause performance issues
- ConcurrentHashMap may give better performance

Static Methods and Synchronized



**Synchronized on an instance method
doesn't synchronize across all instances**

- Problem if static variables are modified

Can use synchronized on static methods

- Class object (E.g. `Foo.class`) is passed to synchronized

Summary




A race condition is caused by threads executing code in an unintended order

- The code where this happens is a critical section

Mutexes protect critical sections

- Obtaining a mutex means only that thread can enter the critical section
- Other threads must wait until it is released

Summary



Java uses the synchronized keyword which is passed an object

- The block becomes a critical section
- Implemented by monitors

Use mutual exclusion to:

- Protect invariants
- Keep object fields consistent
- Prevent simultaneous reads and writes

Summary



AtomicInteger


- Atomic increment and decrement

AtomicBoolean

- 'test-and-set'

Saw Hashtable to prevent concurrent access to a collection

Summary




Synchronized on a code block

- Object passed should be final and not shared elsewhere

Synchronized on a method

- Instance method uses object 'this' refers to
- Static method uses class objects

Summary



Java mutexes are reentrant

- Try to obtain one we're holding, we won't have to block

And now you should be able to answer the following riddle:

How many threads does it
take to change a lightbulb?



How many threads does it
take to change a lightbulb?

All of them... if you don't use
a mutex!

